Forecasting Crude Oil Returns using News Sentiment and Machine Learning *

Tianyu Du †

Wednesday 5th February, 2020

Contents

*Compile Date: 23:44 Wednesday 5th February, 2020

†tianyu.du@mail.utoronto.ca

1 Introduction

2 Data

In order to identify the predictive power of sentiment data on crude oil returns, this study involves three major datasets, a the daily spot price of crude oil at the West Texas Intermediate (WTI) from which returns are computed, ii) a news sentiment dataset from Ravenpack News Analytics (RPNA), and iii) other macroeconomic indicators proxying the overall economic background.

2.1 The West Texas Intermediate (WTI) Crude Oil Dataset

2.2 Summary Statistics of Crude Oil Dataset

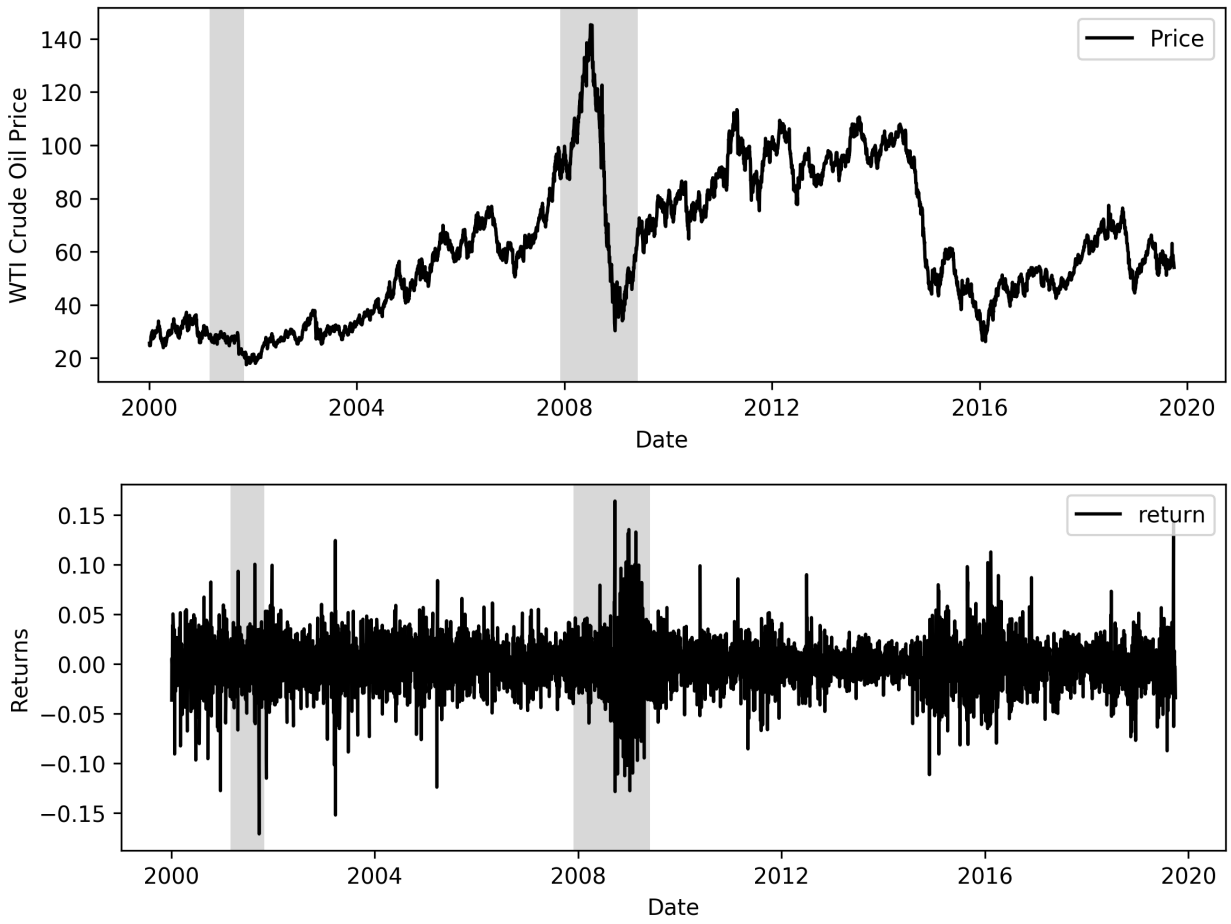

Figure 1: Crude oil prices and returns between January 1, 2000 and September 30, 2019. Shaded areas indicate U.S. recessions.

From the **the table below** we can see that the mean returns for crude oil are around zero year by year. During periods of recessions (March 2001 to November 2001 and December 2007 to June 2009), the data exhibited negative mean returns as well as relatively high standard deviations.

Year	Num. Obs.	Mean	Median	Std.	Min	Max	ACF(1)	ACF(3)	ACF(5)
2000	249	0.000	0.004	0.029	-0.127	0.083	0.012	-0.007	0.126
2001	250	-0.001	-0.001	0.029	-0.171	0.101	0.024	-0.005	-0.037
2002	250	0.002	0.002	0.021	-0.062	0.060	-0.030	-0.008	-0.014
2003	250	0.000	0.002	0.028	-0.152	0.124	-0.133	0.096	-0.097
2004	249	0.001	0.003	0.023	-0.076	0.059	-0.074	0.019	-0.036
2005	251	0.001	0.002	0.022	-0.124	0.084	-0.085	-0.083	-0.109
2006	249	-0.000	0.001	0.018	-0.049	0.062	0.002	0.008	-0.030
2007	252	0.002	0.001	0.019	-0.047	0.055	-0.103	0.000	0.069
2008	253	-0.003	-0.001	0.039	-0.128	0.164	0.008	0.165	-0.259
2009	252	0.002	0.002	0.034	-0.127	0.133	-0.034	0.096	-0.022
2010	252	0.001	0.000	0.019	-0.052	0.099	0.051	-0.071	0.057
2011	252	0.000	0.001	0.022	-0.085	0.086	0.027	-0.003	-0.087
2012	252	-0.000	0.001	0.016	-0.048	0.090	-0.154	0.034	0.120
2013	252	0.000	0.001	0.011	-0.035	0.032	0.045	-0.073	-0.153
2014	252	-0.002	-0.001	0.016	-0.111	0.049	-0.209	0.054	-0.042
2015	252	-0.001	-0.004	0.029	-0.091	0.098	-0.113	-0.106	-0.021
2016	252	0.001	0.000	0.031	-0.080	0.113	0.006	-0.040	0.078
2017	250	0.000	0.003	0.016	-0.056	0.033	-0.017	-0.017	0.076
2018	249	-0.001	0.001	0.020	-0.077	0.073	-0.103	-0.056	0.011
2019	187	0.001	0.001	0.023	-0.087	0.142	-0.090	-0.039	0.123
Total	4955	0.000	0.001	0.024	-0.171	0.164	-0.035	0.021	-0.024

Table 1: Summary Statistics for Crude Oil Returns in each Year. Note that this dataset only include nine months of 2019.

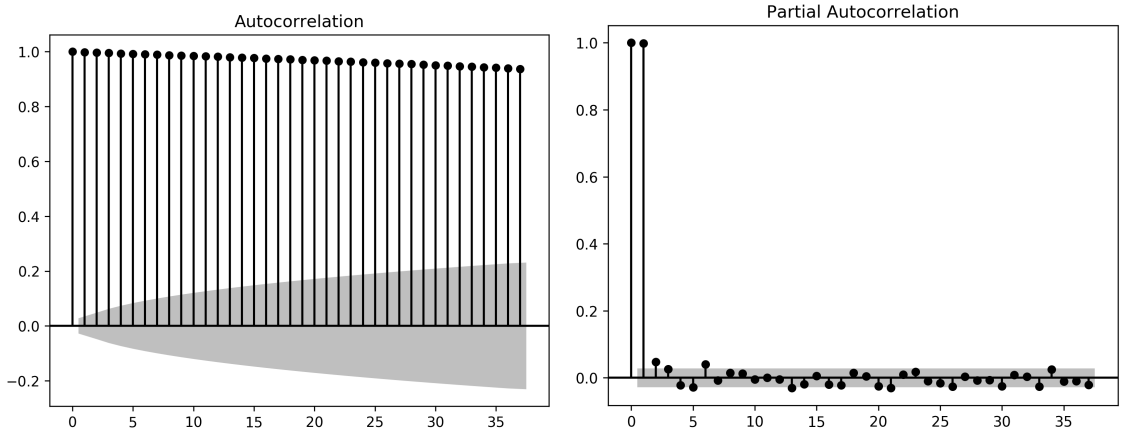

Figure 2: ACF and PACF for Crude Oil Prices (January 2000 to September 2019)

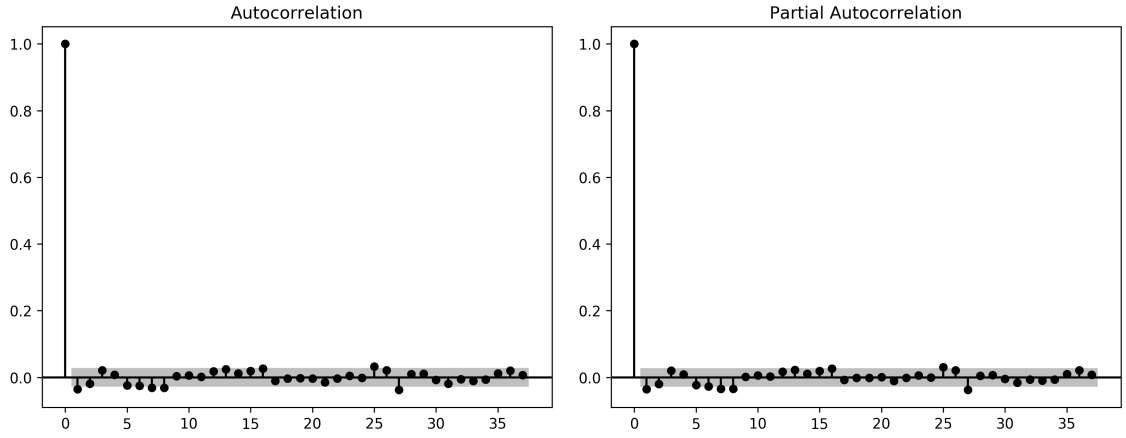

Figure 3: ACF and PACF for Crude Oil Returns (January 2000 to September 2019)

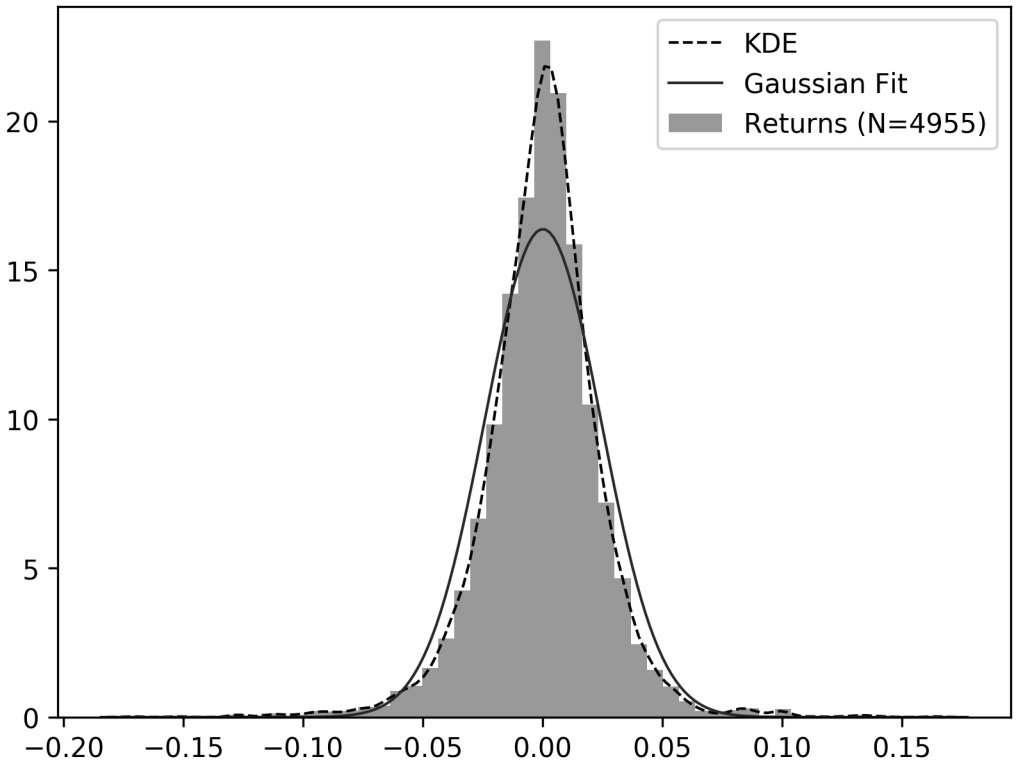

Figure 4: Distribution of Crude Oil Returns (January 2000 to September 2019). The KDE line stands for the kernel density estimation for the empirical distribution. The 'Gaussian Fit' line corresponds to the density function of $\mathcal{N}(\hat{\mu}_{\text{sample}}, \hat{\sigma}_{\text{sample}})$.

2.3 Missing Data in Crude Oil Dataset

This paper calculates crude oil returns on one particular day t by taking the difference in logged prices at t and the previous trading day:

$$r_t := \ln(p_t) - \ln(p_{t-\Delta}) \tag{2.1}$$

where $t - \Delta$ is the last trading day before day t. Different gaps between consecutive trading days, Δ, have salient influence on actual realized returns. Therefore, missing data issue is crucial for this paper.

As mentioned before, the time gap between two observed prices are not equal. For instance, the return on a Monday can be computed by taking difference between the log close price on Monday and the previous Friday, if available. In this case, $\Delta = 3$. If the previous Friday was a holiday without valid price data, r_t will be $\ln(p_{\text{Mon}}) - \ln(p_{\text{Prev Thu}})$, and $\Delta = 4$. According to **the table below**, 33 days are in this case.

Day of the week	Num. Days.	Num. Trading Days	Δ=1	2	3	4	5
Monday	1031	927	0	0	883	33	11
Tuesday	1030	1018	921	0	0	97	0
Wednesday	1030	1022	1011	5	0	0	6
Thursday	1030	1002	994	8	0	0	0
Friday	1030	986	969	17	0	0	0
Saturday	1030	0	0	0	0	0	0
Sunday	1030	0	0	0	0	0	0
Total	7211	4955	3895	30	883	130	17

Table 2: The values of Δ used to calculate returns. This table only include trading days, but the first day with price observation in this dataset was dropped because it did not have a previous trading day, so return on this day cannot by computed using our definition.

As mentioned before, the oil price dataset does not have any prices over weekends. **The table below** reports dates that are most frequently associated with a missing data over the span of 20 years. The pool of days with missing data is pretty consistent overtime, the market is always closed on January 1, July 4 (Independence Day) and December 25 (Christmas). The group of dates in late November are responsible for missing data on Thanksgiving holiday.

Date	Counts (all)	Counts (excl. weekends)
July 4	20	16
January 1	20	14
December 25	19	14
July 3	10	5
November 23	10	5
November 24	10	4
November 25	10	3
November 22	9	4
November 26	9	3

Table 3: Dates most frequently associates with missing data. Data on January 1, July 4, and December 25 are missing ever year. Because the entire dataset ranges from January 3, 2000 to September 30, 2019, missing data problems on December 25 are only reported 19 times.

2.4 Day of the Week Effect in Crude Oil Dataset

2.4.1 Difference in Returns across the Week

Gibbons and Hess' work examined returns on stocks from S&P 500, Dow Jones 30, and Treasury Bills. They found strong negative mean returns on Monday compared with other weekdays. The seasonality persisted even after taking market adjustment measures, such as using mean-adjusted returns instead (Gibbons & Hess, 1981). Analysis in my paper unveils a similar daily seasonality presents in crude oil returns as well. **Panels in the figure below** demonstrate the empirical distributions of returns on each day of the week. We can see that Mondays and Wednesdays have relatively larger variances, which again matches Gibbons and Hess' observations.

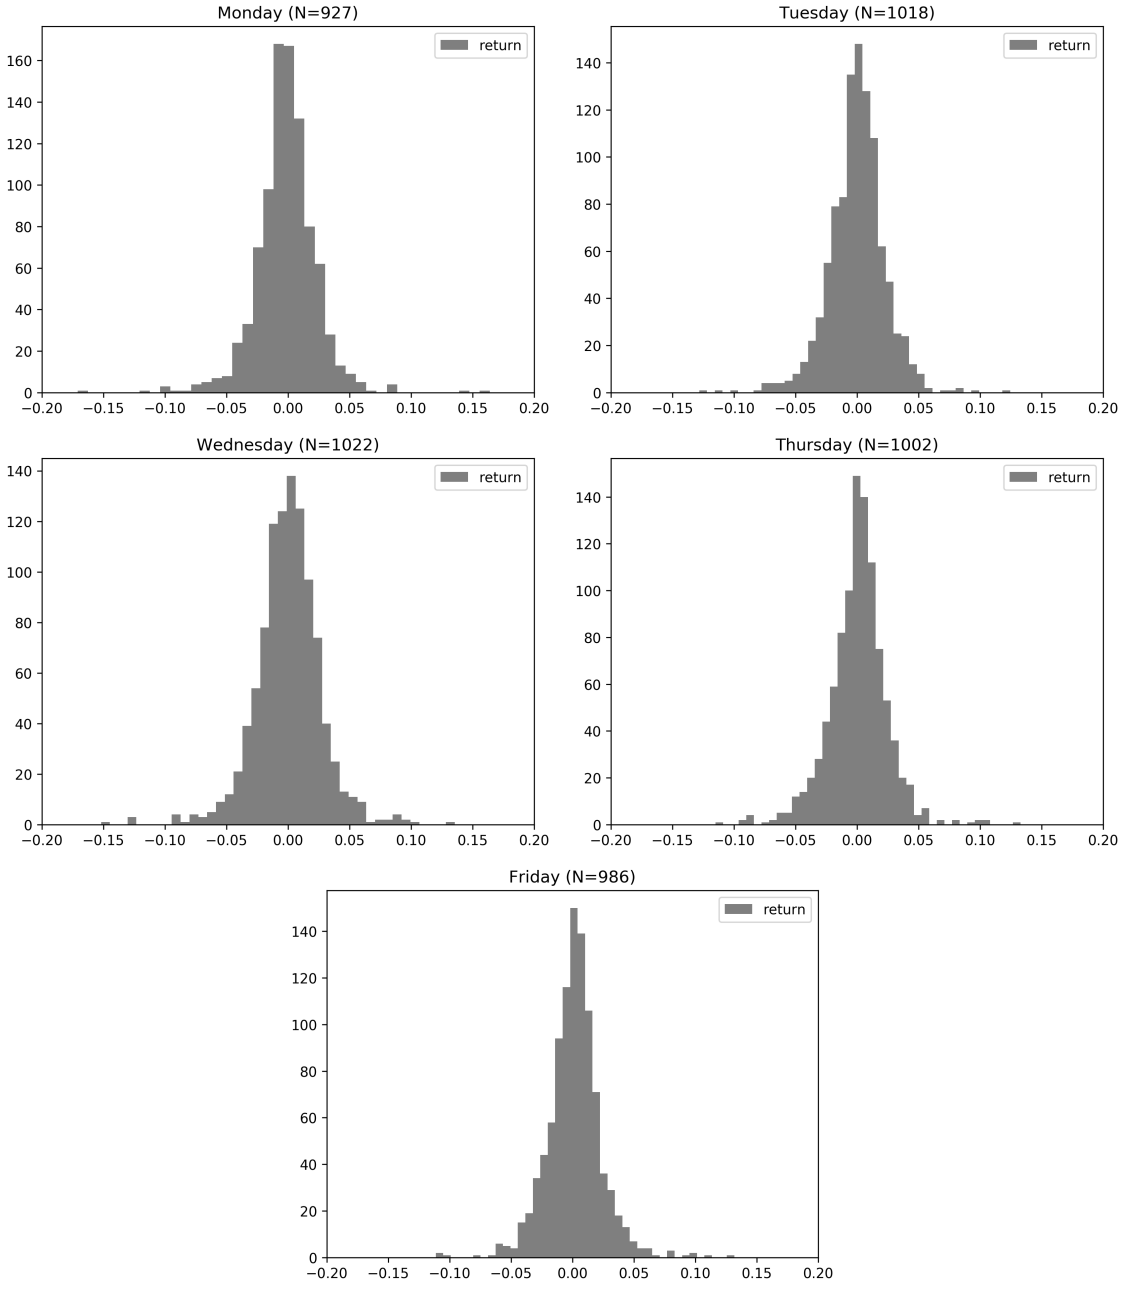

Figure 5: Crude oil returns on each weekday. Weekend data are not available in the daily dataset provided by U.S. Energy Information Administration (EIA). Ns within parentheses in figure titles denote the number of observations. See appendix for distributions of crude oil prices.

The **two tables** below provide summary statistics for prices and returns on each day. It turns out that Monday is the only weekday with a mean return significantly less than zero.

Day of the week	Num. Obs.	Mean	Std.	3^{rd} Moment
Monday	927	62.072	26.493	7081.163
Tuesday	1019	61.828	26.317	6895.638
Wednesday	1022	61.810	26.398	7049.810
Thursday	1002	62.005	26.431	6955.555
Friday	986	62.079	26.247	6676.566
Total	4956			

Table 4: Summary statistics of crude oil prices on each day of week

Day of the week	Num. Obs.	Mean (P-Value)	Std.	3^{rd} Moment
Monday	927	**-0.002 (0.049)**	0.025	-0.0000019
Tuesday	1018	-0.000 (0.900)	0.023	-0.0000031
Wednesday	1022	0.000 (0.884)	0.027	-0.0000054
Thursday	1002	0.001 (0.361)	0.024	-0.0000006
Friday	986	**0.002 (0.0311)**	0.023	0.0000021
Total	4955			

Table 5: Summary statistics of crude oil returns on each day of week. The first day (January 1, 2000) of the oil price dataset was Saturday, and the observation on the following Monday (January 3) was missing. Hence, the return on Tuesday (January 4) could not be computed because it was the first trading day in this dataset, and there are only 1018 Tuesdays in the dataset of returns. A value of -0.000 indicates a negative value with magnitude less than 0.0005. P-values are calculated in a two-tailed t-test with $\mu_0 = 0$. Bold fonts indicate statistically significance at level $\alpha = 0.05$.

2.4.2 Kolmogorov-Smirnov test for Distributional Similarities

Smirnov developed a non-parametric method of testing the equality between two continuous distributions, with CDFs $F(x)$ and $G(x)$ respectively, (Smirnov, 1939). Refer to Hodges' work for a detailed review on the Kolmogorov-Smirnov test (Hodges, 1958). I am using the two-tailed version of Kolmogorov-Smirnov test to check whether distributions of two different days are similar. Given two datasets, take returns on Mondays and Tuesdays for example, the null hypothesis says those two datasets are drawn from the same distribution, and the alternative says they are from different distributions [1]. Firstly, the Kolmogorov–Smirnov test constructs the empirical CDFs $F_{Mon,927}(x)$

[1]Different alternative hypotheses can be used in Kolmogorov–Smirnov test: i) $H_1 : F(x) \geq G(x)$, ii) $H_1 : F(x) \leq G(x)$, and iii) $H_1 : F(x) \neq G(x)$. This paper is using the third (two-tailed) alternative hypothesis.

and $F_{Tue,1018}(x)$ from the dataset. Then, the Kolmogorov–Smirnov statistic measures the maximum discrepancy between two distribution functions, which is

$$D := \sup_x |F_{Mon,927}(x) - F_{Tue,1018}(x)| \in [0, 1] \tag{2.2}$$

A smaller D-statistic implies stronger distributional similarity between two distributions. For instance, when $F_{Mon,927}(x)$ and $F_{Tue,1018}(x)$ are exactly the same, the D-statistic is zero. In contrast, let $X = 0$ and $Y = 1$ be two deterministic random variables, in this case, $D_{X,Y} = 1$.
The test rejects H_0 at a significance level of α if

$$D > \sqrt{-\frac{1}{2} \ln \frac{\alpha}{2}} \sqrt{\frac{n+m}{nm}} \tag{2.3}$$

where m and n denote sizes of two datasets.

D-Statistic (P-Value)	Monday	Tuesday	Wednesday	Thursday	Friday
Monday	0.000 (1.000)	**0.061 (0.048)**	**0.065 (0.030)**	**0.092 (0.001)**	**0.092 (0.001)**
Tuesday		0.000 (1.000)	**0.044 (0.260)**	0.036 (0.505)	0.044 (0.264)
Wednesday			0.000 (1.000)	0.053 (0.114)	**0.073 (0.009)**
Thursday				0.000 (1.000)	0.025 (0.900)
Friday					0.000 (1.000)

Table 6: The Kolmogorov-Smirnov D-Statistic for all pairs of distributions. Bold font indicates the null hypothesis is rejected at a significance level of 0.05, which implies discrepancy in distributions.

The table above presents the Kolmogorov-Smirnov D-Statistic for distributions of every pairs of days. At a significance level of 0.05, we can see that Mondays follow a distribution significantly different from distributions of other weekdays follow. Because the dataset does not contain weekend data, returns on Mondays is always computed using the difference between log prices on Monday and the previous Friday (Thursday if Friday is not a trading day and so on). Therefore, returns associated with Mondays pick the weekend effect. In fact, the distribution of returns on Mondays (over weekends) is the only one with negative mean among distributions of all five days.

2.5 News and Sentiment Datasets

TODO: *Need revision: how to describe the dataset efficiently.*

The event sentiment dataset from RavenPack News Analytics (RPNA) tracks and analyzes all information of companies, organizations, countries, commodities, and currencies from four major sources: Dow Jones Newswires, Wall Street Journal, Barron's and MarketWatch.

The dataset covers events from January 1, 2000, to September 30, 2019. RavenPack records the exact date and coordinated universal time (UTC) when each news is published.

For each piece of news, the dataset links it to a unique entity name attribute. To filter out noise data less relevant to crude oil returns, this paper selects the subset of news with crude oil topic. There are 106, 960 entries from the original dataset left, lead to 15 events per day on average. In the figure below, panel A presents a distribution of ESS for all news related to crude oil in the time span of 20 years and panel B shows all distributions of events within each year.

Moreover, the dataset categorizes each event following the RavenPack taxonomy.

(i) topic;

(ii) group;

(iii) type;

(iv) sub-type;

(v) property;

(vi) category: fine details.

Figure 6: Ravenpack taxonomy TODO: *Add examples of each level.* TODO: *Add definitions of each level.*

To proxy the potential economic impact upon news arrival and afterwards, Ravenpack assigns each piece of news an Event Sentiment Score (ESS) between 0 and 100 using an algorithm combines results from surveying financial experts and pattern matching. An ESS of 100 indicates extreme positive short-term positive financial or economic impact. In contrast, a 0 ESS score indicates extreme negative impact. And a ESS of 50 indicates exact neutral news. From this point, scores are normalized by subtracting 50, so that the sign of normalized ESS matches the nature of news, and a zero score represents a neutral news. **The histogram below** plots the distribution of normalized ESS for all news about crude oil. It turns out that only a small portion of news is purely neutral (i.e., with zero ESS) TODO: *(Probably move this part to the 'classification' section.)*

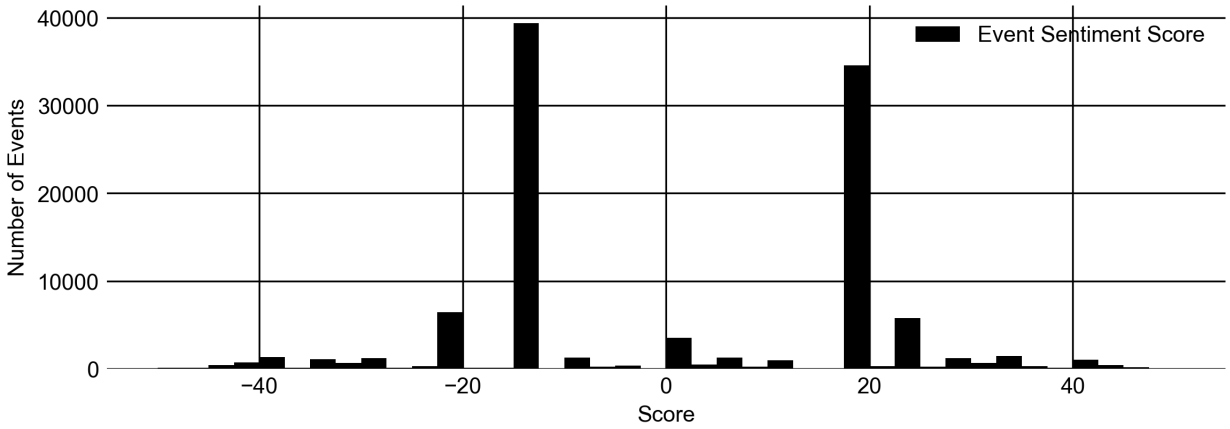

Figure 7: Distribution of Event Sentiment Scores of all 106,960 news items.

It is worth mentioning that ESS measures the potential impact on the topic of this news. For example, a civil unrest in a middle east country is often considered as news forerun negative economic impact, especially for the country itself. However, such news is in general associated with positive ESS scores because the expected negative supply shocks carried by these news are typically positively correlated crude oil prices and returns. **Tables below** present a list of categories frequently associated with positive and negatives news. From these two table we can see that the majority of themes of positive news would impact crude oil prices and returns positively.

Category	Number of positive news
commodity-price-gain	22,893
commodity-futures-gain	11,648
supply-decrease-commodity	5,845
imports-up	2,705
commodity-buy-target	1,171
demand-increase-commodity	1,070
exports-down	1,020
other 28 categories	3,014
all positive news	49,366

Table 7: Most frequent categories of positive news. Only categories with frequency greater than 1,000 are shown in this table.

Category	Number of negative news
commodity-price-loss	26,475
commodity-futures-loss	12,818
supply-increase-commodity	6,629
imports-down	2,017
exports-up	1308
resource-discovery-commodity	1,179
technical-view-bearish	1,172
other 24 categories	2,517
all negative news	54,115

Table 8: Most frequent categories of positive news. Only categories with frequency greater than 1,000 are shown in this table.

2.6 Classifying News Type

2.7 Case Studies of Events

2.7.1 Positive Spike on November 30, 2016

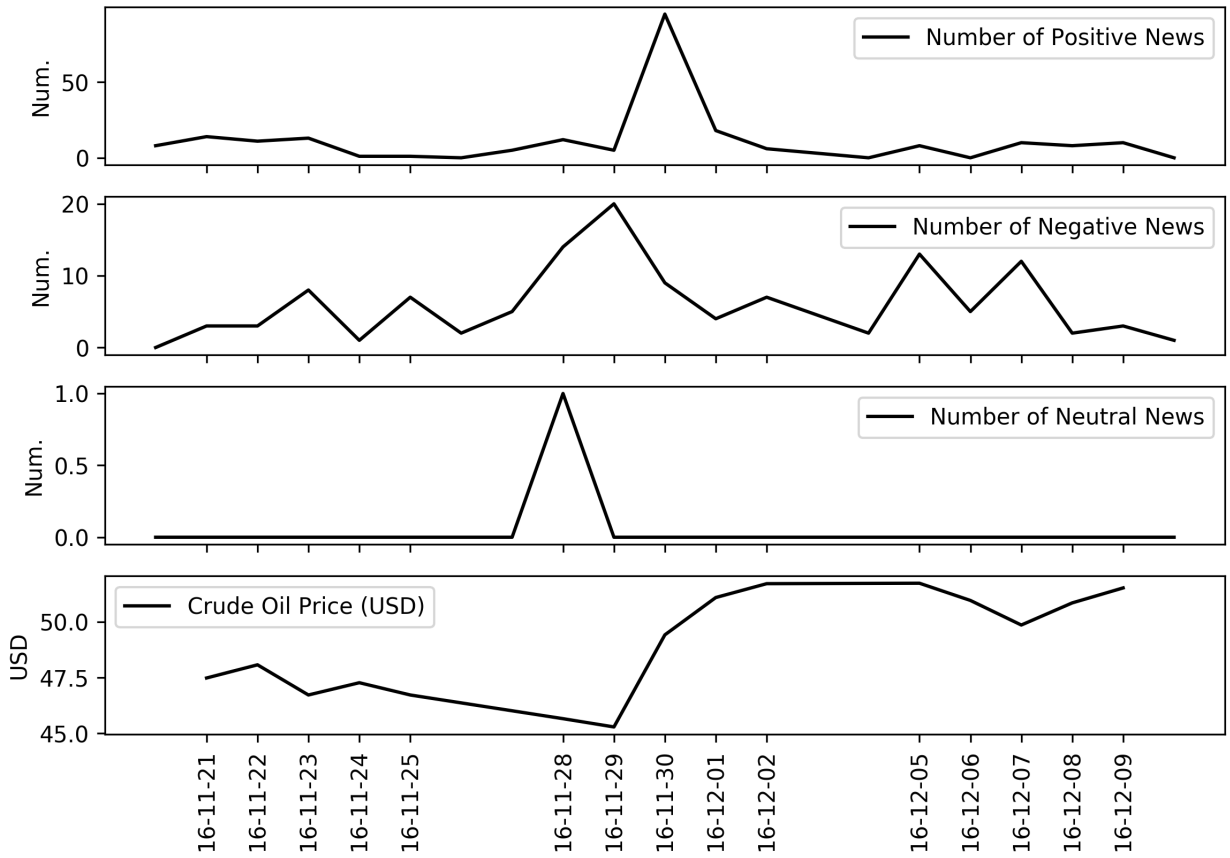

Figure 8:

2.7.2 Negative Spike on December 6, 2018

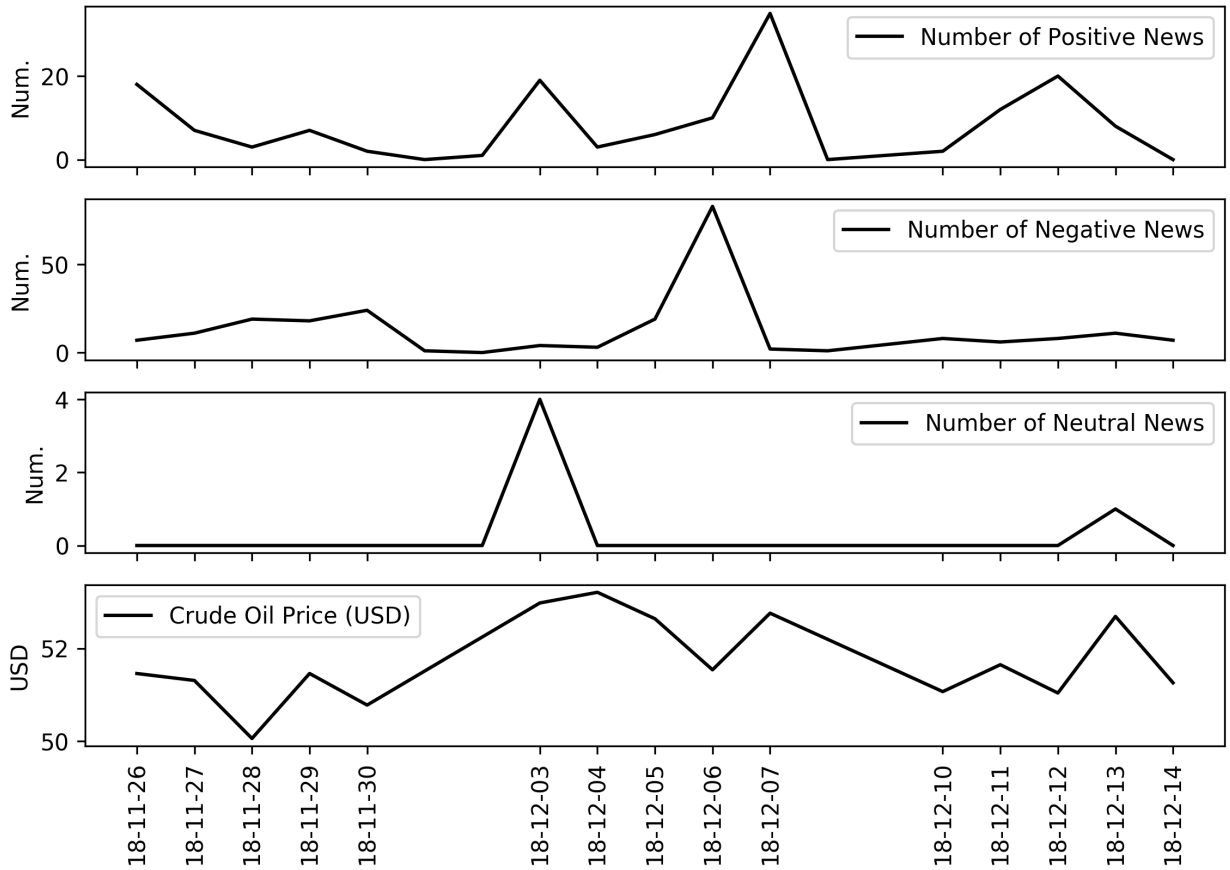

Figure 9:

2.7.3 Positive Spike on June. 12 - 13, 2019

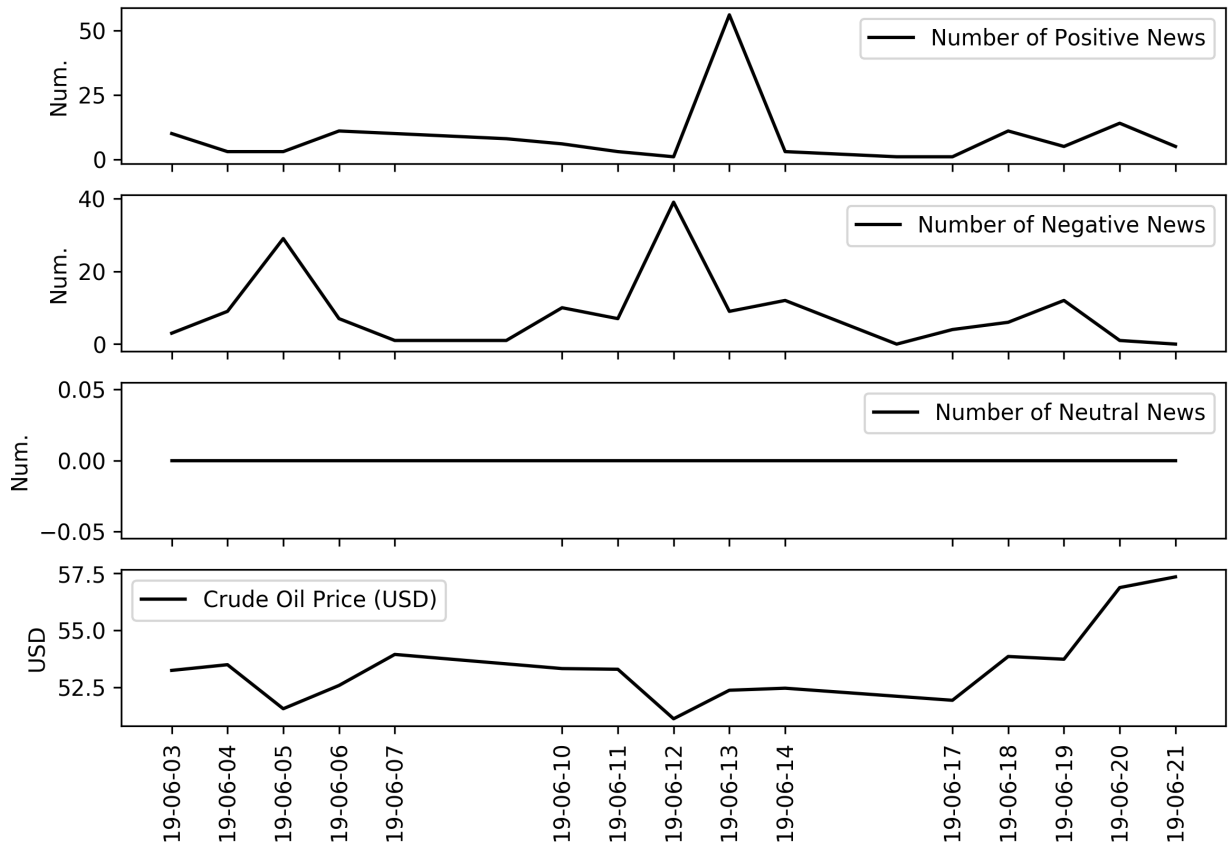

Figure 10:

3 Models

4 Experiments

References

Gibbons, M. R., & Hess, P. (1981). Day of the Week Effects and Asset Returns. *The Journal of Business*, *54*(4), 579. doi: 10.1086/296147

Hodges, J. L. (1958). The significance probability of the smirnov two-sample test. *Arkiv för Matematik*, *3*(5), 469–486. doi: 10.1007/bf02589501

Smirnov, N. (1939). On the estimation of the discrepancy between empirical curves of distribution for two independent samples. *Bulletin Moscow University*, *2*, 3–16.

5 Appendix: Supplementary Summary Statistics for Datasets

Year	Num. Obs.	Mean	Median	Std.	Min	Max	ACF(1)	ACF(3)	ACF(5)
2000	250	30.379	30.270	2.966	23.910	37.220	0.946	0.838	0.740
2001	250	25.983	27.185	3.560	17.500	32.210	0.973	0.924	0.880
2002	250	26.185	26.700	3.208	18.020	32.680	0.976	0.925	0.870
2003	250	31.075	30.770	2.624	25.250	37.960	0.943	0.864	0.764
2004	249	41.506	40.700	5.775	32.490	56.370	0.982	0.949	0.915
2005	251	56.637	57.330	6.252	42.160	69.910	0.969	0.917	0.876
2006	249	66.055	65.650	5.586	55.900	77.050	0.975	0.929	0.893
2007	252	72.341	69.735	12.853	50.510	99.160	0.986	0.956	0.924
2008	253	99.672	104.830	28.563	30.280	145.310	0.986	0.958	0.926
2009	252	61.950	67.025	13.361	34.030	81.030	0.985	0.959	0.928
2010	252	79.476	79.735	5.242	64.780	91.480	0.953	0.853	0.759
2011	252	94.881	95.790	8.063	75.400	113.390	0.968	0.900	0.828
2012	252	94.053	92.605	7.713	77.720	109.390	0.979	0.946	0.914
2013	252	97.983	96.325	5.451	86.650	110.620	0.977	0.927	0.881
2014	252	93.172	97.850	13.519	53.450	107.950	0.978	0.936	0.895
2015	252	48.657	47.870	6.814	34.550	61.360	0.972	0.928	0.888
2016	252	43.294	45.080	6.727	26.190	54.010	0.978	0.932	0.893
2017	250	50.800	50.385	3.914	42.480	60.460	0.968	0.905	0.846
2018	249	65.227	66.380	6.517	44.480	77.410	0.961	0.888	0.812
2019	187	57.037	56.580	3.986	46.310	66.240	0.924	0.811	0.732
Total	4956	61.956	58.915	26.376	17.500	145.310	0.998	0.995	0.992

Table 9: Summary Statistics for Crude Oil Prices. Note that this dataset only include nine months of 2019.

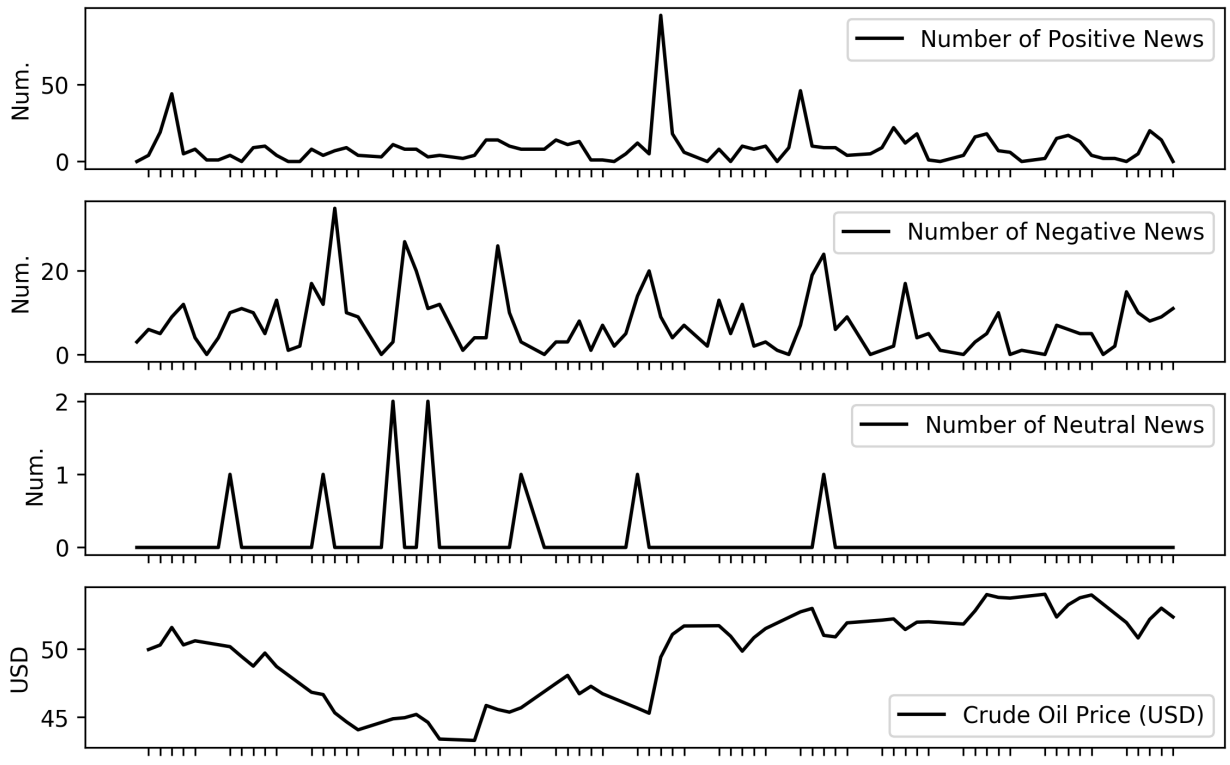
Num.
50
0
Number of Positive News
Num.
20
0
Number of Negative News
Num.
2
1
0
Number of Neutral News
USD
50
45
Crude Oil Price (USD)

Figure 11:

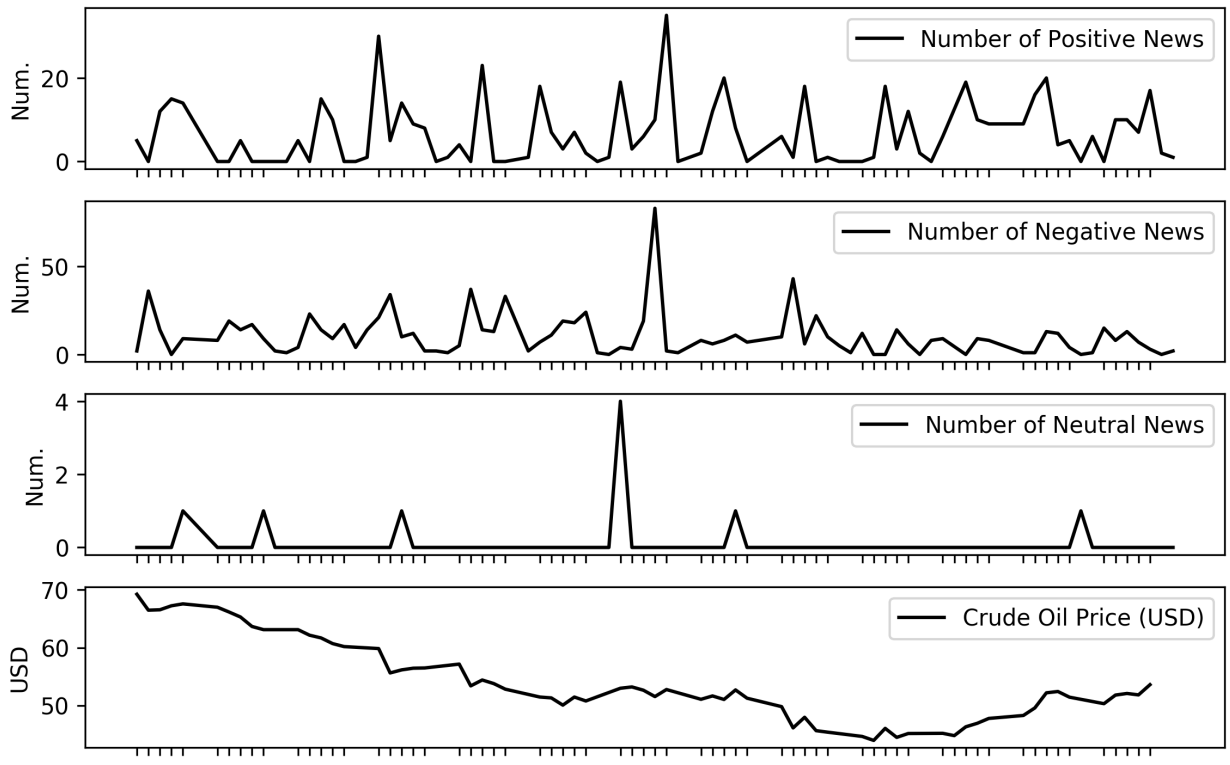
Number of Positive News
Num.
20
0
Number of Negative News
Num.
50
0
Number of Neutral News
Num.
4
2
0
Crude Oil Price (USD)
USD
70
60
50

Figure 12:

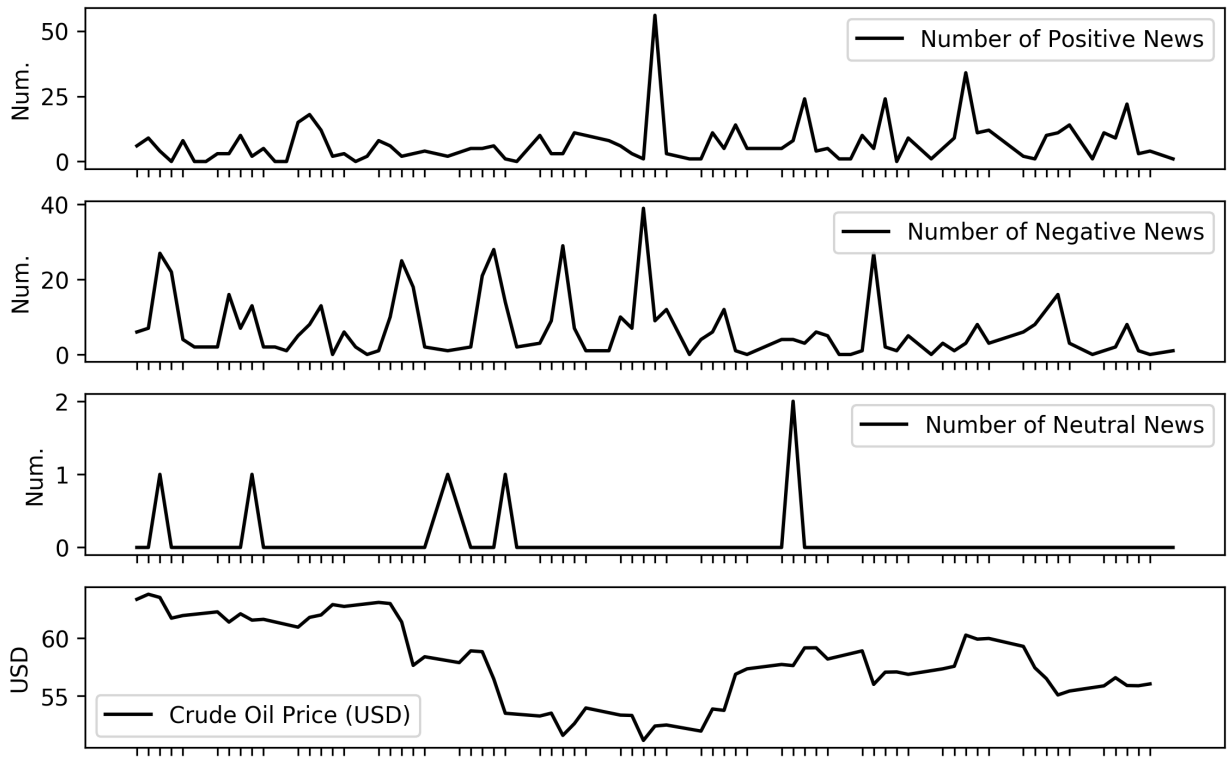
Number of Positive News
Num.
50
25
0
Number of Negative News
Num.
40
20
0
Number of Neutral News
Num.
2
1
0
Crude Oil Price (USD)
USD
60
55

Figure 13: